

# Smarthome Device Efficiency Prediction

```
In [2]: # Edit all the Markdown cells below with the appropriate information
# Run all cells, containing your code
# Save this Jupyter with the outputs of your executed cells
# PS: Save again the notebook with this outcome.
# PSPS: Don't forget to include the dataset in your submission
```

**Team:**

- Danny Nguyen

**Course:** CISC 43 – BIG DATA (Spring, 2025)

## Problem Statement

- This project uses a dataset titled "smart\_home\_device\_usage\_data.csv" found on Kaggle. The objective will be to predict whether a device is efficient or inefficient. I will aim to do this using supervised classification machine learning algorithms.
- Keywords:** efficiency classification, smart home device

## Required packages

- Add instructions to install the required packages

```
In [7]: ## Your code begins here
import pandas as pd
import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

## Methodology

Description:

This dataset captures smart home device usage metrics, offering insights into user behavior, device efficiency, and preferences. It includes data on device types, usage patterns, energy consumption, malfunction incidents, and user satisfaction metrics.

Features:

- UserID: Unique identifier for each user.
- DeviceType: Type of smart home device (e.g., Lights, Thermostat).
- UsageHoursPerDay: Average hours per day the device is used.
- EnergyConsumption: Daily energy consumption of the device (kWh).
- UserPreferences: User preference for device usage (0 - Low, 1 - High).
- MalfunctionIncidents: Number of malfunction incidents reported.
- DeviceAgeMonths: Age of the device in months.
- SmartHomeEfficiency (Target Variable): Efficiency status of the smart home device (0 - Inefficient, 1 - Efficient).

1. Explain your big data methodology

I will first try to predict whether a device is considered efficient or inefficient using the K-NN methodology

2. Introduce the topics you used in your project

- Model 1
  - KNN
- Model 2
  - Linear Regression

Your code starts here

## Initial Exploratory Data Analysis

```
In [11]: ## Loading dataset
df = pd.read_csv("smart_home_device_usage_data.csv")
```

```
In [12]: df.describe()
```

	UserID	UsageHoursPerDay	EnergyConsumption	UserPreferences	MalfunctionIncidents	DeviceAgeMonths	SmartHomeEfficiency
count	5403.000000	5403.000000	5403.000000	5403.000000	5403.000000	5403.000000	5403.000000
mean	2702.000000	12.052992	5.054302	0.511753	2.066445	30.312234	0.376643
std	1559.856083	6.714961	2.878941	0.499908	1.423291	16.990525	0.484589
min	1.000000	0.501241	0.101562	0.000000	0.000000	1.000000	0.000000
25%	1351.500000	6.297871	2.524968	0.000000	1.000000	15.000000	0.000000
50%	2702.000000	11.903768	5.007047	1.000000	2.000000	30.000000	0.000000
75%	4052.500000	17.791751	7.611912	1.000000	3.000000	45.000000	1.000000
max	5403.000000	23.987326	9.998071	1.000000	4.000000	59.000000	1.000000

```
In [13]: df.info()
## Inspect data types
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5403 entries, 0 to 5402
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   UserID                 5403 non-null   int64
1   DeviceType             5403 non-null   object
2   UsageHoursPerDay       5403 non-null   float64
3   EnergyConsumption      5403 non-null   float64
4   UserPreferences        5403 non-null   int64
5   MalfunctionIncidents   5403 non-null   int64
6   DeviceAgeMonths        5403 non-null   int64
7   SmartHomeEfficiency    5403 non-null   int64
dtypes: float64(2), int64(5), object(1)
memory usage: 337.8+ KB

```

```

In [14]: ## From above we can see that there are no nulls, but I will make sure.
        ## There are no missing values either
        print(df.isnull().sum())
        ## Looks Like this dataset is fairly clean!

```

```

UserID                0
DeviceType            0
UsageHoursPerDay      0
EnergyConsumption     0
UserPreferences       0
MalfunctionIncidents  0
DeviceAgeMonths       0
SmartHomeEfficiency   0
dtype: int64

```

```

In [15]: ## But I know that the datatypes are wrong based on the description of Features in the original dataset page
        ## For example, DeviceType should be Categorical.

```

```

df['DeviceType'] = df['DeviceType'].astype('category')
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5403 entries, 0 to 5402
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   UserID                 5403 non-null   int64
1   DeviceType             5403 non-null   category
2   UsageHoursPerDay       5403 non-null   float64
3   EnergyConsumption      5403 non-null   float64
4   UserPreferences        5403 non-null   int64
5   MalfunctionIncidents   5403 non-null   int64
6   DeviceAgeMonths        5403 non-null   int64
7   SmartHomeEfficiency    5403 non-null   int64
dtypes: category(1), float64(2), int64(5)
memory usage: 301.1 KB

```

```

In [16]: df.head()

```

```

Out[16]:
   UserID  DeviceType  UsageHoursPerDay  EnergyConsumption  UserPreferences  MalfunctionIncidents  DeviceAgeMonths  SmartHomeEfficiency
0      1    Smart Speaker      15.307188          1.961607              1              4              36              1
1      2      Camera      19.973343          8.610689              1              0              29              1
2      3  Security System      18.911535          2.651777              1              0              20              1
3      4      Camera       7.011127          2.341653              0              3              15              0
4      5      Camera      22.610684          4.859069              1              3              36              1

```

```

In [17]: df.tail()

```

```

Out[17]:
   UserID  DeviceType  UsageHoursPerDay  EnergyConsumption  UserPreferences  MalfunctionIncidents  DeviceAgeMonths  SmartHomeEfficiency
5398  5399   Thermostat      4.556314          5.871764              1              0              28              0
5399  5400     Lights      0.561856          1.555992              1              4              24              0
5400  5401  Smart Speaker      11.096236          7.677779              0              0              42              0
5401  5402  Security System      8.782169          7.467929              0              2              28              1
5402  5403   Thermostat      13.540381          9.043076              0              0              30              0

```

```

In [18]: ## I want to see the correlation between variables

```


```

numeric_df = df.select_dtypes(include=['number'])
correlation_matrix = numeric_df.corr()

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f", linewidths=0.5)
plt.title("Feature Correlation Matrix")
plt.show()

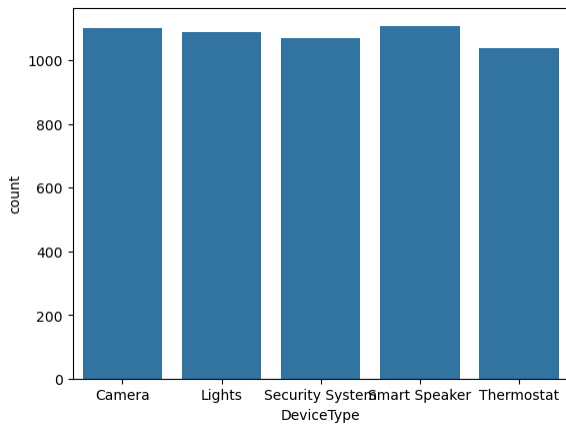
## Interestingly, nearly all variables have little correlation with each other, with the exception of UserPreferences and SmartHomeEfficiency
## This makes sense, because a person would likely to use an Efficient device with High preference.

```



```
In [22]: sns.countplot(data=df, x='DeviceType')
```

```
Out[22]: <Axes: xlabel='DeviceType', ylabel='count'>
```



## K-Nearest Neighbors Analysis

- I want to categorize whether a device is efficient or inefficient.
- The dataset already includes a column for this target variable, "SmartHomeEfficiency", but I want to see if we can apply machine learning methods using other features to predict it.

```
In [30]: # Note that I had converted "DeviceType" into a categorical variable. To do KNN, I need to convert this to a number through encoding
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df["DeviceType"] = le.fit_transform(df["DeviceType"])

## I need to define the features (X) and the target (y)
# For the features, I will remove UserID because that is just a unique identifier, and I will remove the target variable
X = df.drop(['UserID', 'SmartHomeEfficiency'], axis=1)
y = df['SmartHomeEfficiency']

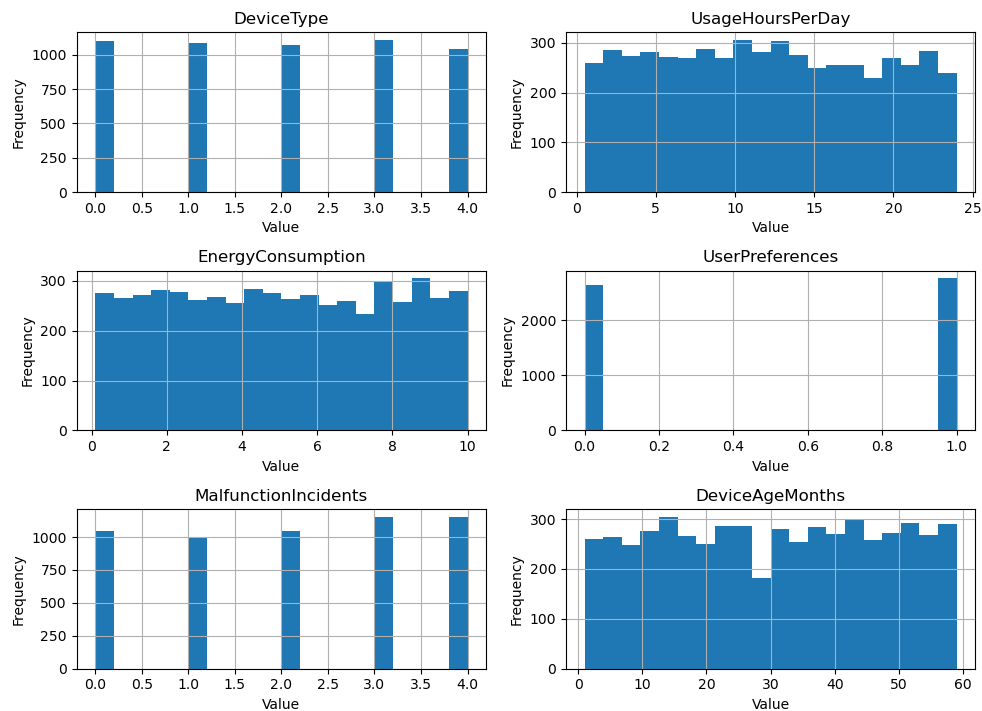
# Splitting the dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [33]: # Plot before scaling
axes = X.hist(figsize=(10, 8), bins=20)

# Iterate over each subplot and set the titles
for ax in axes.flatten():
    ax.set_xlabel('Value')
    ax.set_ylabel('Frequency')

plt.suptitle('Histograms of Features Before Scaling')
plt.tight_layout(rect=[0, 0, 1, 0.95]) # Adjust the padding to make room for the suptitle
plt.show()
```

## Histograms of Features Before Scaling



```
In [34]: # Feature Scaling
scaler = StandardScaler()
scaler.fit(X_train)

X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [35]: # Create KNN classifier
k = 5 # Number of neighbors
knn = KNeighborsClassifier(n_neighbors=k)

# Train the classifier
knn.fit(X_train, y_train)

# Predict on the test data
y_pred = knn.predict(X_test)

# Evaluate the model
print(confusion_matrix(y_test, y_pred))
```

```
[[637  41]
 [ 44 359]]
```

- The confusion matrix above shows the summary of prediction results. The results of this model are mostly accurate--637 were correctly negative, 359 were correctly positive, while only 41 were incorrectly positive, and only 44 were incorrectly negative.

```
In [40]: # Classification report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.94	0.94	0.94	678
1	0.90	0.89	0.89	403
accuracy			0.92	1081
macro avg	0.92	0.92	0.92	1081
weighted avg	0.92	0.92	0.92	1081

## Logistic Regression Analysis

```
In [43]: ## Logistic Regression would also be appropriate for this because it is binomial output
from sklearn.linear_model import LogisticRegression
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_predlog = logreg.predict(X_test)

print(confusion_matrix(y_test, y_predlog))
```

```
[[608  70]
 [ 65 338]]
```

```
In [45]: print(classification_report(y_test, y_predlog))
```

	precision	recall	f1-score	support
0	0.90	0.90	0.90	678
1	0.83	0.84	0.83	403
accuracy			0.88	1081
macro avg	0.87	0.87	0.87	1081
weighted avg	0.88	0.88	0.88	1081

Conclusions

- The results of our KNN model are also fairly good. The scores are fairly high, indicating that our model is mostly accurate.
- Results of logistic regression model are also not bad, though it performs slightly worse than using KNN.
- The Confusion Matrix for both methods show that the models predicted correctly in most cases.
- These results are not surprising because this is a simple, synthetic dataset that was likely designed for beginner students to practice on.

References

- CISD 43 Course Files and Examples
- <https://medium.com/aiskunks/categorical-data-encoding-techniques-d6296697a40f>

Credits

- I looked at code from our CISD 43 course "KNN\_Classification\_Example 1", "KNN\_Classification\_Example 2", and "KNN\_Classification\_Example 3" as reference for code to set up the K-NN analysis
- I also reference code from our CISD 43 examples relating to MongoDB, Module 12 NoSQL MongoDB Example 1 and Module 12 NoSQL Tutorial and MongoDB Example 2 for examples of query code used together with MongoDB.
- Thanks to the author Krishnakanth Naik Jarapala at this site, <https://medium.com/aiskunks/categorical-data-encoding-techniques-d6296697a40f>, I was able to find the Python code needed to encode categorical datatypes

Mongo DB Connection

```
In [52]: import pymongo
from pymongo import MongoClient
# Connect to MongoDB
client = MongoClient("mongodb://localhost:27017/")

In [53]: db = client['smarthomedeviceefficiency_db']
collection = db['deviceefficiency']

In [54]: # Retrieve the first three records
retrieve = collection.find().limit(3)
# Print the records
for record in retrieve:
    print(record)

{'_id': ObjectId('6842295f14b42ab475966aeb'), 'UserID': 1, 'DeviceType': 'Smart Speaker', 'UsageHoursPerDay': 15.30718848124909, 'EnergyConsumption': 1.9616068166289793, 'UserPreference
s': 1, 'MalfunctionIncidents': 4, 'DeviceAgeMonths': 36, 'SmartHomeEfficiency': 1}
{'_id': ObjectId('6842295f14b42ab475966aec'), 'UserID': 2, 'DeviceType': 'Camera', 'UsageHoursPerDay': 19.9733432937798, 'EnergyConsumption': 8.610688921898104, 'UserPreferences': 1, 'M
alfunctionIncidents': 0, 'DeviceAgeMonths': 29, 'SmartHomeEfficiency': 1}
{'_id': ObjectId('6842295f14b42ab475966aed'), 'UserID': 3, 'DeviceType': 'Security System', 'UsageHoursPerDay': 18.91153466115779, 'EnergyConsumption': 2.651776634718286, 'UserPreferenc
es': 1, 'MalfunctionIncidents': 0, 'DeviceAgeMonths': 20, 'SmartHomeEfficiency': 1}

In [58]: allrecords = collection.find()
mongodf = pd.DataFrame(list(allrecords))
print(mongodf)

   _id  UserID  DeviceType  UsageHoursPerDay  \
0  6842295f14b42ab475966aeb      1  Smart Speaker      15.307188
1  6842295f14b42ab475966aec      2    Camera      19.973343
2  6842295f14b42ab475966aed      3  Security System      18.911535
3  6842295f14b42ab475966aee      4    Camera       7.011127
4  6842295f14b42ab475966aef      5    Camera      22.610684
...    ...      ...      ...      ...
5398 6842295f14b42ab475968001  5399    Thermostat      4.556314
5399 6842295f14b42ab475968002  5400      Lights      0.561856
5400 6842295f14b42ab475968003  5401  Smart Speaker      11.096236
5401 6842295f14b42ab475968004  5402  Security System      8.782169
5402 6842295f14b42ab475968005  5403    Thermostat      13.540381

   EnergyConsumption  UserPreferences  MalfunctionIncidents  \
0           1.961607              1              4
1           8.610689              1              0
2           2.651777              1              0
3           2.341653              0              3
4           4.859069              1              3
...           ...              ...              ...
5398          5.871764              1              0
5399          1.555992              1              4
5400          7.677779              0              0
5401          7.467929              0              2
5402          9.043076              0              0

   DeviceAgeMonths  SmartHomeEfficiency
0              36              1
1              29              1
2              20              1
3              15              0
4              36              1
...           ...              ...
5398              28              0
5399              24              0
5400              42              0
5401              28              1
5402              30              0

[5403 rows x 9 columns]

In [60]: allrecords_count = collection.count_documents({})
print(f"Total # of Records: {allrecords_count}")

Total # of Records: 5403

In [62]: cameradevices = collection.count_documents({"DeviceType": "Camera"})
print("Total number of Camera Records:", cameradevices)

Total number of Camera Records: 1101

In [64]: # End of Project

In [ ]:
```